

```

[1] import sys
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import rcParams
import seaborn as sns
from scipy.stats import zscore
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split, cross_val_score

```

```

[2] # code in this cell from:
# https://stackoverflow.com/questions/27934885/how-to-hide-code-from-cells-in-ipython-notebook-visualized-with-nbviewer
from IPython.display import HTML

HTML('''<script>
code_show=true;
function code_toggle() {
  if (code_show){
    $('div.input').hide();
  } else {
    $('div.input').show();
  }
  code_show = !code_show
}
$( document ).ready(code_toggle);
</script>
<form action="javascript:code_toggle()"><input type="submit" value="Click here to display/hide the code."></form>''')

```

Click here to display/hide the code.

```

# switch to seaborn default stylistic parameters
sns.set()
sns.set_context('notebook')

```

```

df = pd.read_csv("https://raw.githubusercontent.com/grbruns/cst383/master/kuiper-2008-cars.csv")
df.drop(['Model', 'Trim'], inplace=True, axis=1)

```

Are there any NA values in the data set?

```

df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 804 entries, 0 to 803
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Price       804 non-null   float64
1   Mileage     804 non-null   int64
2   Make        804 non-null   object
3   Type        804 non-null   object
4   Cylinder    804 non-null   int64
5   Liter       804 non-null   float64
6   Doors       804 non-null   int64
7   Cruise      804 non-null   int64
8   Sound       804 non-null   int64
9   Leather     804 non-null   int64
dtypes: float64(2), int64(6), object(2)
memory usage: 62.9+ KB

```

✓
5s

```
# 1 Produce a grid of scatterplots using only Price, Mileage, and Liter  
# Note: use a semicolon after your last plot statement to suppress  
# the non-graphical output.
```

```
# YOUR CODE HERE
```

```
sns.pairplot(df[['Price', 'Mileage', 'Liter']], diag_kws={'bins': 10})
```

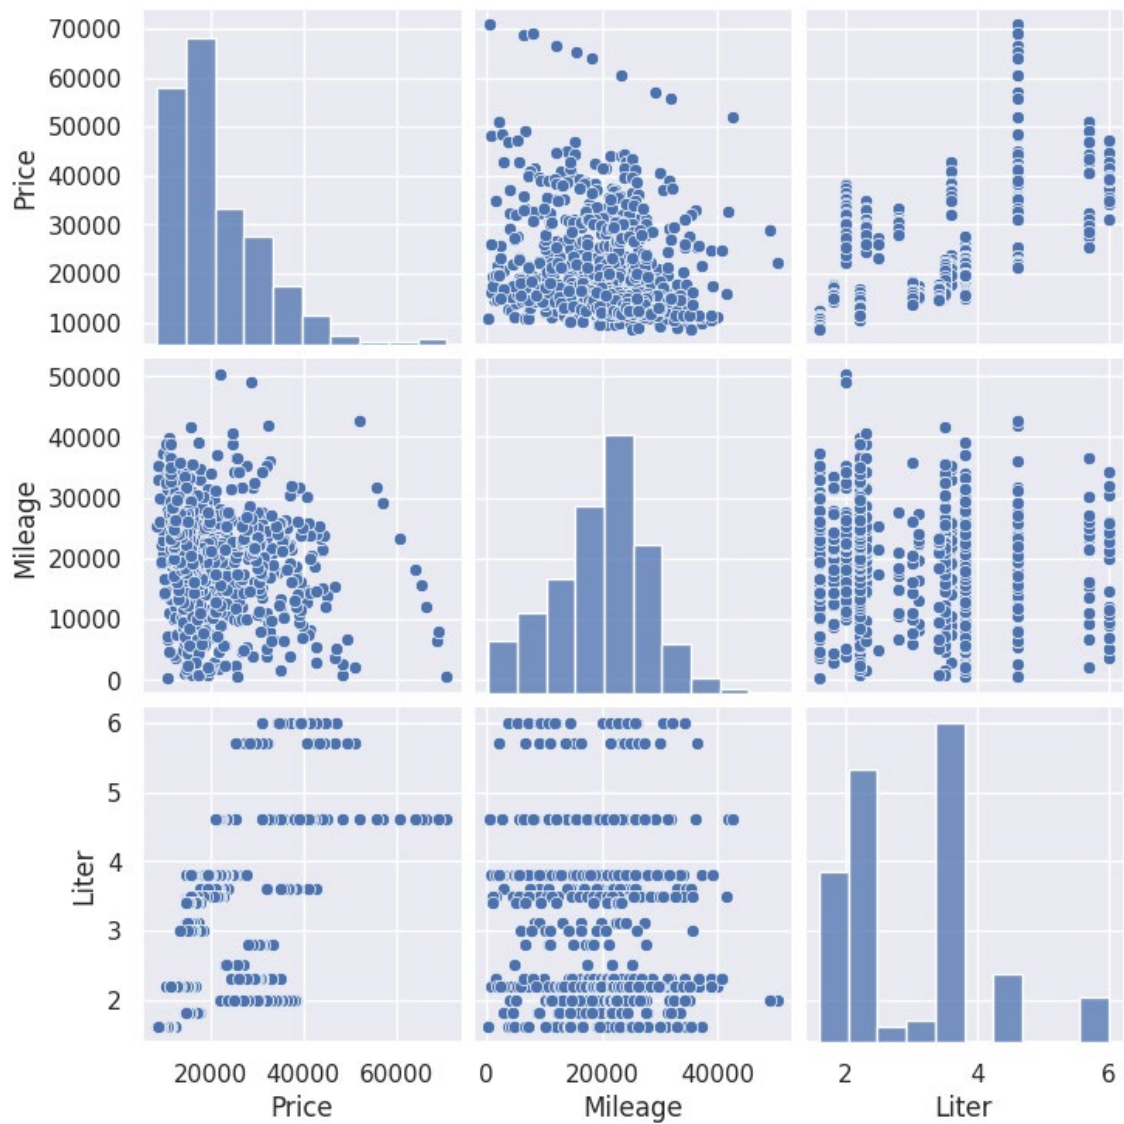


```
<seaborn.axisgrid.PairGrid at 0x7c79d35cc940>
```

✓
5s



```
<seaborn.axisgrid.PairGrid at 0x7c79d35cc940>
```



✓
0s

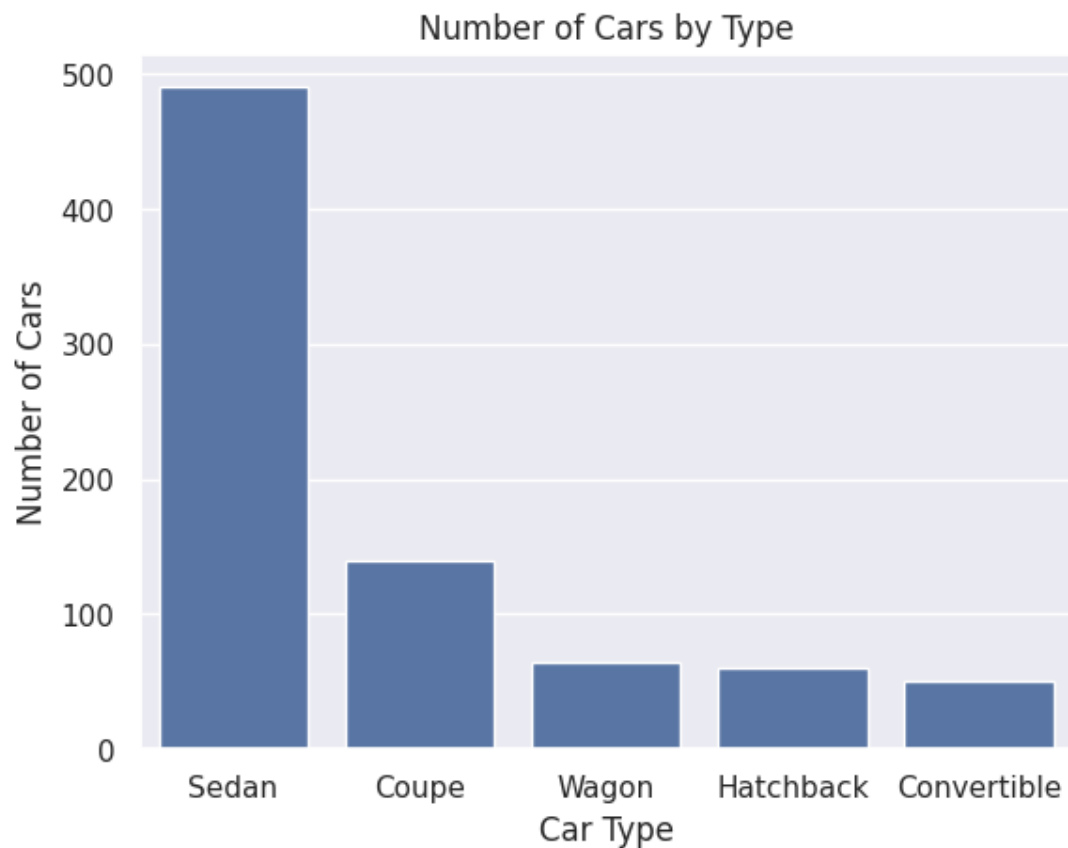


```
# YOUR CODE HERE
car_counts = df['Type'].value_counts()

# Create a bar plot
plt.figure(figsize=(10, 8))
sns.barplot(x=car_counts.index, y=car_counts.values)
plt.title('Number of Cars by Type')
plt.xlabel('Car Type')
plt.ylabel('Number of Cars')
```



Text(0, 0.5, 'Number of Cars')



Let's build a model to predict a used car's price from its mileage.

✓
0s

▶ @@ 3 Build a linear model using Mileage as the single predictor variable,
and Price as the target variable. Fit the model using the entire data set.
Use the LinearRegression class and use variable 'reg' for your LinearRegression object.

```
# YOUR CODE HERE
X = df[['Mileage']]
y = df['Price']

reg = LinearRegression()

reg.fit(X, y)
```



```
LinearRegression
LinearRegression()
```


✓
0s



```
#@ 4 Create a scatterplot of Price by Mileage (Price on y axis),  
# then superimpose your linear model on it (as a line). Do not use  
# Seaborn's regplot for this -- use two plotting statements.  
  
# YOUR CODE HERE  
plt.scatter(df['Mileage'], df['Price'], label='Data')  
  
plt.plot(df['Mileage'], reg.predict(X), color='red', linewidth=2, label='Linear Model')  
  
plt.title('Price vs Mileage with Linear Model')  
plt.xlabel('Mileage')  
plt.ylabel('Price')
```



Text(0, 0.5, 'Price')



✓
0s



```
#@ 5 Print the coefficients and R-squared value of the model.  
# Hint: to get the R-squared value you can use the score()  
# method of LinearRegression.
```

```
# YOUR CODE HERE  
coefficients = reg.coef_  
intercept = reg.intercept_  
r_squared = reg.score(X, y)  
  
print(f"intercept: {intercept:.2f}")  
print(f"coefficients for Mileage: {coefficients[0]:.2f}")  
print(f"r-squared value: {r_squared:.2f}")
```



```
intercept: 24764.56  
coefficients for Mileage: -0.17  
r-squared value: 0.02
```

✓
0s



```
#@ 6 Create a new linear model for Price using predictors Mileage,  
# cruise, and Leather. Assign your model to variable reg2.
```

```
# YOUR CODE HERE  
X = df[['Mileage', 'Cruise', 'Leather']]  
  
y = df['Price']  
  
reg2 = LinearRegression()  
  
reg2.fit(X, y)
```



```
LinearRegression  
LinearRegression()
```

✓
0s



```
## 7 Print the coefficients (including intercept) for the new model.  
# Hint: (Use print()'{:.2f}'.format(x)) to print a value  
# x with 2 digits after the decimal point.)
```

```
# YOUR CODE HERE
```

```
coefficients2 = reg2.coef_  
intercept2 = reg2.intercept_
```

```
print(f"Intercept: {intercept2:.2f}")  
print("Coefficient:")  
print(f"  Mileage: {coefficients2[0]:.2f}")  
print(f"  for Cruise: {coefficients2[1]:.2f}")  
print(f"  for Leather: {coefficients2[2]:.2f}")
```



```
Intercept: 14297.18  
Coefficient:  
  Mileage: -0.19  
  for Cruise: 10256.12  
  for Leather: 4175.58
```

✓
0s



```
## 8 Using your model, compute the predicted price for  
# a car with a mileage of 20,000 and cruise but no leather.  
# Hint: create a matrix to predict() that contains only one row.  
# Print the predicted price.
```

```
# YOUR CODE HERE
```

```
car_mil_20k_c = [[20000, 1, 0]]
```

```
predicted_price = reg2.predict(car_mil_20k_c)
```

```
print(f"{predicted_price[0]:.2f}")
```




```
20827.73  
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names  
  warnings.warn()
```

✓
0s



```
## 9 Define a function named 'plot_actual_predicted' to plot predicted vs actual values.  
# It should take as input a NumPy array of actual values, a NumPy array of  
# predicted values, and a plot title. The two arrays can be assumed to be the same length.  
# I used linewidth=2 and linestyle='dashed' for the actual-predicted line.  
#  
# Hint: when plotting, first plot the scatter plot and then plot the line that  
# shows when predicted=actual. You need only two points to plot the line where  
# actual = predicted, and the points should have the form (a,a), (b,b).  
# To determine what a and b should be, you may want to compute 1) the minimum  
# value of actual and predicted, and 2) the maximum value of actual and predicted.
```

```
def plot_actual_predicted(actual, predicted, title):  
    # YOUR CODE HERE  
    plt.scatter(actual, predicted, edgecolor='white', label='Predicted vs Actual')  
  
    min_val = min(np.min(actual), np.min(predicted))  
    max_val = max(np.max(actual), np.max(predicted))  
  
    plt.plot([min_val, max_val], [min_val, max_val], color='black', linestyle='dashed', linewidth=2, label='Actual = Predicted')  
  
    plt.xlabel('Actual Values')  
    plt.ylabel('Predicted Values')  
    plt.title(title)  
    plt.legend()
```



```
[17] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
```

Now let's fit our same model using only training data.

```
#@ 10 Create another linear model (again building a model to predict Price from
# Mileage, Cruise, and Leather). Call your new model reg3.
# However, this time fit the model using the training data.

# YOUR CODE HERE
reg3 = LinearRegression()

reg3.fit(X_train, y_train)
```

LinearRegression
LinearRegression()


✓
0s




```
# Use the training data.
```

```
# YOUR CODE HERE
```

```
plot_actual_predicted(y_train, reg3.predict(X_train), 'Predicted vs Actual (Training Set)')
```



```
#@ 12 Plot the actual and predicted values using plot_actual_predicted().  
# This time use the test data.  
  
# YOUR CODE HERE  
plot_actual_predicted(y_test, reg3.predict(X_test), 'Predicted vs Actual (Test Set)')
```



✓ 0s

```
#@ 13 Print the root mean squared error on the test data. This is the  
# square root of the average squared error. Write your own code  
# to compute the RMSE; don't use a library function.
```

```
# YOUR CODE HERE  
y_test_pred = reg3.predict(X_test)  
  
squared_differences = (y_test - y_test_pred) ** 2  
  
mean_squared_error = np.mean(squared_differences)  
  
rmse = np.sqrt(mean_squared_error)  
  
print(f"RMSE: {rmse:.2f}")
```

↔ RMSE: 8517.23

✓
15

```
#@ 14 Create a new model reg4 that is like reg3, but adds 'Cylinder'
# as a new predictor. Do a train/test split (with random_state = 0), and fit your model to
# the training data.

# YOUR CODE HERE
X_new = df[['Mileage', 'Cruise', 'Leather', 'Cylinder']]

#X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
X_train, X_test, y_train, y_test = train_test_split(X_new, y, test_size=0.25, random_state=0)

reg4 = LinearRegression()

reg4.fit(X_train, y_train)
```



LinearRegression
LinearRegression()

✓
05

```
# Print the RMSE of your new model on the test data,
# and R-squared value of your new model (on the training data)

y_test_pred_reg4 = reg4.predict(X_test)
sqr_diff = (y_test - y_test_pred_reg4) ** 2
mean_squared_error = np.mean(sqr_diff)
rmse = np.sqrt(mean_squared_error)

#print('RMSE of reg4: {:.2f}'.format(np.sqrt(mean_squared_error(y_test, reg4.predict(X_test))))) #TypeError: 'numpy.float64' object is not callable
print('RMSE of reg4: {:.2f}'.format(rmse))
print('r-squared value of reg4: {:.4f}'.format(reg4.score(X_train, y_train)))
```




RMSE of reg4: 7810.77
r-squared value of reg4: 0.4489

✓
0s



```
#@ 15 Plot the actual and predicted values using plot_actual_predicted().  
# Use test data for your predictions.  
  
# YOUR CODE HERE  
plot_actual_predicted(y_test, y_test_pred_reg4, 'Predicted by Actual (reg4)')
```



Some of the predictions are still really bad. Would scaling the data help us make better predictions?

✓
0s



```
#@ 16 Using predictors Mileage, Cruise, Leather, and Cylinder, make  
# NumPy arrays X and y, where y contains the values in column Price.  
# then scale all columns of array X using scipy.stats.zscore.  
# Use X_s as the name of the scaled version of X.  
  
# YOUR CODE HERE  
X = df[['Mileage', 'Cruise', 'Leather', 'Cylinder']].values  
y = df['Price'].values  
X_s = zscore(X)
```

```

X_train, X_test, y_train, y_test = train_test_split(X_s, y, test_size=0.25, random_state=0)

regs = LinearRegression()
regs.fit(X_train, y_train)

print('r-squared value of regs: {:.4f}'.format(regs.score(X_train, y_train)))

```

→ r-squared value of regs: 0.4489

✓ [27] # so that, if the categorical variable has n different unique
0s # values, then only n-1 dummy variables will be used.

```

# YOUR CODE HERE
df = pd.get_dummies(df, drop_first=True)

```

A check to ensure that we now have only numeric variables.

✓ df.info()
0s

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 804 entries, 0 to 803
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Price                 804 non-null   float64
1   Mileage               804 non-null   int64
2   Cylinder              804 non-null   int64
3   Liter                 804 non-null   float64
4   Doors                 804 non-null   int64
5   Cruise                804 non-null   int64
6   Sound                 804 non-null   int64
7   Leather               804 non-null   int64
8   Make_Cadillac         804 non-null   bool
9   Make_Chevrolet        804 non-null   bool
10  Make_Pontiac           804 non-null   bool
11  Make_SAAB              804 non-null   bool
12  Make_Saturn            804 non-null   bool
13  Type_Coupe             804 non-null   bool
14  Type_Hatchback         804 non-null   bool
15  Type_Sedan             804 non-null   bool
16  Type_Wagon             804 non-null   bool
dtypes: bool(9), float64(2), int64(6)
memory usage: 57.4 KB

```

✓
0s

```
# YOUR CODE HERE
y = df['Price'].values
X = df.drop('Price', axis=1).values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)

reg5 = LinearRegression()
reg5.fit(X_train, y_train)

print(f"Intercept: {reg5.intercept_:.2f}")
print("Coefficients:")
for feature, coef in zip(df.drop('Price', axis=1).columns, reg5.coef_):
    print(f" {feature}: {coef:.2f}")
```

```
↔ Intercept: 34393.88
Coefficients:
Mileage: -0.19
Cylinder: -1217.73
Liter: 5674.20
Doors: -5338.01
Cruise: 149.11
Sound: 150.82
Leather: 112.58
Make_Cadillac: 15638.25
Make_Chevrolet: -1616.20
Make_Pontiac: -1831.79
Make_SAAB: 10623.51
Make_Saturn: -1345.82
Type_Coupe: -12563.40
Type_Hatchback: -2209.53
Type_Sedan: -2221.69
Type_Wagon: 1762.21
```

▶ #@ 19 Print the r-squared value for your model based on the training data
and also print the RMSE based on the test data.

```
# YOUR CODE HERE
r_squared_train = reg5.score(X_train, y_train)
print(f"R-squared: {r_squared_train:.2f}")

y_test_pred_reg5 = reg5.predict(X_test)
sqr_diff = (y_test - y_test_pred_reg5) ** 2
mean_squared_error = np.mean(sqr_diff)
rmse = np.sqrt(mean_squared_error)

print(f"RMSE: {rmse:.2f}")
```

⇒ R-squared: 0.94
RMSE: 2685.22

✓
0s

▶ #@ 20 From your NumPy array X create an extended data set X_poly using
PolynomialFeatures with degree=2. Assign your PolynomialFeatures
object to variable pf.

```
# YOUR CODE HERE
pf = PolynomialFeatures(degree=2)

X_poly = pf.fit_transform(X)

print(X.shape)
```

⇒ (804, 16)

✓
0s

▶ X_poly.shape

⇒ (804, 153)

Let's go for broke and build a model using *all* features. What is the RMSE on the test data for such a model?

✓
0s



```
## 21 Create a model reg6 using all of these features. First  
# create training and test sets (using X_poly and y), then  
# train a linear model on the training data, and then compute  
# the RMSE on the test data.
```

```
# YOUR CODE HERE
```

```
X_train_poly, X_test_poly, y_train, y_test = train_test_split(X_poly, y, test_size=0.25, random_state=0)
```

```
reg6 = LinearRegression()  
reg6.fit(X_train_poly, y_train)
```



LinearRegression

LinearRegression()



```
#@ 22 Using the ideas in the last cell, compute the RMSE for each individual
# feature using 5-fold cross validation. Save the index and RMSE associated
# with the best feature as the two variables i_min and rmse_min.

# YOUR CODE HERE
#i_min = None
#rmse_min = float('inf')
#
#for i in range(X.shape[1]):
#    X_single_feature = X[:, i].reshape(-1, 1)
#    X_train, X_test, y_train, y_test = train_test_split(X_single_feature, y, test_size=0.3, random_state=42)
#
#    model = LinearRegression()
#    model.fit(X_train, y_train)
#
#    y_test_pred = model.predict(X_test)
#    sqr_diff = (y_test - y_test_pred) ** 2
#    mean_squared_error_val = np.mean(sqr_diff)
#    rmse = np.sqrt(mean_squared_error_val)
#
#    print(f"Feature index {i} RMSE: {rmse:.2f}")
#
#    if rmse < rmse_min:
#        rmse_min = rmse
#        i_min = i

#print(f"num features: {i_min}")
#print(f"Minimum RMSE: {rmse_min:.2f}")

def get_rmse(model, X, y, n_splits=5):
    n_samples = len(y)
    indices = np.arange(n_samples)
    np.random.shuffle(indices)
    fold_sizes = np.full(n_splits, n_samples // n_splits, dtype=int)
    fold_sizes[:n_samples % n_splits] += 1
    current = 0
```

```

rmse_values = []

for fold_size in fold_sizes:
    start, stop = current, current + fold_size
    test_indices = indices[start:stop]
    train_indices = np.concatenate((indices[:start], indices[stop:]))
    current = stop

    X_train, X_test = X[train_indices], X[test_indices]
    y_train, y_test = y[train_indices], y[test_indices]

    model.fit(X_train, y_train)
    y_test_pred = model.predict(X_test)
    sqr_diff = (y_test - y_test_pred) ** 2
    mean_squared_error_val = np.mean(sqr_diff)
    rmse = np.sqrt(mean_squared_error_val)
    rmse_values.append(rmse)

return np.mean(rmse_values)

i_min = None
rmse_min = float('inf')

for i in range(X.shape[1]):
    X_single_feature = X[:, i].reshape(-1, 1)
    model = LinearRegression()
    rmse = get_rmse(model, X_single_feature, y)
    print(f"Feature index {i} RMSE: {rmse:.2f}")

    if rmse < rmse_min:
        rmse_min = rmse
        i_min = i

```

Feature index 0 RMSE: 9786.02
Feature index 1 RMSE: 8083.76
Feature index 2 RMSE: 8187.80
Feature index 3 RMSE: 9771.89
Feature index 4 RMSE: 8902.98
Feature index 5 RMSE: 9787.09
Feature index 6 RMSE: 9751.17
Feature index 7 RMSE: 7436.61
Feature index 8 RMSE: 9013.57
Feature index 9 RMSE: 9779.46
Feature index 10 RMSE: 9280.63
Feature index 11 RMSE: 9644.46
Feature index 12 RMSE: 9740.64
Feature index 13 RMSE: 9655.63
Feature index 14 RMSE: 9886.52
Feature index 15 RMSE: 9865.82

✓
0s

```
# make a version of the training data with just feature 0
X_0 = X_train[:,[0]]

# compute negated mean square error scores using 5-fold cross validation
scores = cross_val_score(LinearRegression(), X_0, y_train, scoring='neg_mean_squared_error', cv=5)

# work out the average root mean squared error. We need to
# first negate the scores, because they are negative MSE, not MSE.
rmse = np.sqrt(-scores.mean())

print('RMSE for feature 0 only: {:.2f}'.format(rmse))
```

RMSE for feature 0 only: 9882.58

✓
5s

```
#print('best feature: {}, best RMSE: {:.2f}'.format(pf.get_feature_names()[i_min], rmse_min))
print('best feature: {}, best RMSE: {:.2f}'.format(pf.get_feature_names_out()[i_min], rmse_min))
```

best feature: x6, best RMSE: 7436.61

✓
1s



```
##@ 23 Find the 10 features that give the lowest RMSE by using
# forward search. Important: find the single best feature, then
# find one more feature that is the best *when combined with
# the single best feature*, and continue, always looking for
# the single best feature combined with all previously-selected
# features.
#
# I've filled in some of the code for you. 'remaining' is a list
# of the features to be considered when finding the next best
# feature. It is initialized to all the features. 'selected' is
# a list of the features that have been chosen in a round of finding the
# next best feature. It is initialized to the empty list.
#
# Don't forget to have include all selected features when looking
# for the next best feature.

remaining = list(range(X_train.shape[1]))
selected = []
n = 10
while len(selected) < n:
    # find the single features that works best in conjunction
    # with the already selected features
    rmse_min = 1e7
    for i in remaining:
        # YOUR CODE HERE
        X_selected = np.concatenate([X_train[:, selected], X_train[:, [i]]], axis=1)
        model = LinearRegression()
        model.fit(X_selected, y_train)
        rmse = get_rmse(model, X_selected, y_train)
        if rmse < rmse_min:
            rmse_min = rmse
            i_min2 = i

    remaining.remove(i_min2)
    selected.append(i_min2)
    print('num features: {}; rmse: {:.2f}'.format(len(selected), rmse_min))
```



```
num features: 1; rmse: 7359.10
num features: 2; rmse: 5995.09
num features: 3; rmse: 3990.66
num features: 4; rmse: 3636.67
num features: 5; rmse: 3456.31
num features: 6; rmse: 3357.37
num features: 7; rmse: 2692.68
num features: 8; rmse: 2544.02
num features: 9; rmse: 2520.73
num features: 10; rmse: 2533.02
```

How does the test RMSE of the model created using the 10 features found with forward feature selection model that uses all the features?

✓
0s



```
## 24 Print the RMSE of the best 10 features on the test set
```

```
# YOUR CODE HERE
```

```
X_train_selected = X_train[:, selected]
```

```
X_test_selected = X_test[:, selected]
```

```
reg_selected = LinearRegression()
```

```
reg_selected.fit(X_train_selected, y_train)
```

```
y_test_pred_selected = reg_selected.predict(X_test_selected)
```

```
rmse_selected = get_rmse(reg_selected, X_test_selected, y_test)
```

```
print("RMSE of the model with the best 10 features:", rmse_selected)
```



```
RMSE of the model with the best 10 features: 2848.202821302465
```